Causal Geometry of an Alg-Zoo RNN: A Mechanistic Case Study

Integrating Causal Ablation, Temporal Dynamics, and Topological Analysis

Kishore Kumar Mariappan, M.Sc.

February 24, 2026

Abstract

A central challenge in mechanistic interpretability is scaling analysis from microscopic weight inspection to macroscopic algorithmic understanding. While sparse circuit discovery has succeeded on minimal networks, it failed on the 432-parameter "Alg-Zoo" RNN ($M_{16,10}$) trained on the `2nd-argmax` task. We hypothesize that this failure stems from the assumption that algorithms must be sparse. Using a novel **Topological Circuit Audit (TCA)** framework, we demonstrate that the network's computation is driven by a distributed, continuous "Logic Ring." We derive a mechanistic error bound of **3.44%** from the manifold's geometric resolution limit—perfectly matching empirical performance—and causally verify the ring's necessity via a controlled topological pinch test.

1 Introduction: The Need for Geometric Diagnostics

In 2026, the Alignment Research Center (ARC) released "Alg-Zoo," a suite of small models designed to benchmark interpretability techniques [1]. While manual sparse circuit analysis [2] succeeded on a 10-parameter RNN, the 432-parameter variant ($M_{16,10}$) remained opaque.

We propose that $M_{16,10}$ utilizes a **distributed geometric representation** rather than a sparse boolean circuit. To reverse-engineer it, we developed the Topological Circuit Audit (TCA), a 5-phase pipeline that bridges the gap between structural ablation, dynamical systems theory, and causal topology.

2 Phase 0: The Attentional Focus

Before analyzing internal circuitry, we must determine what the model attends to. We computed the gradient of the decision margin with respect to the input sequence across 2,000 samples, aligning gradients by the relative rank of input values (Figure 1).

The analysis reveals a precise heuristic: the model ignores the bottom 7 ranks and focuses entirely on the top three. The negative spike at Rank 7 indicates that the 3rd-largest value acts as a repulsive constraint in the decision geometry, isolating the geometric gap (g) as the primary axis of difficulty.

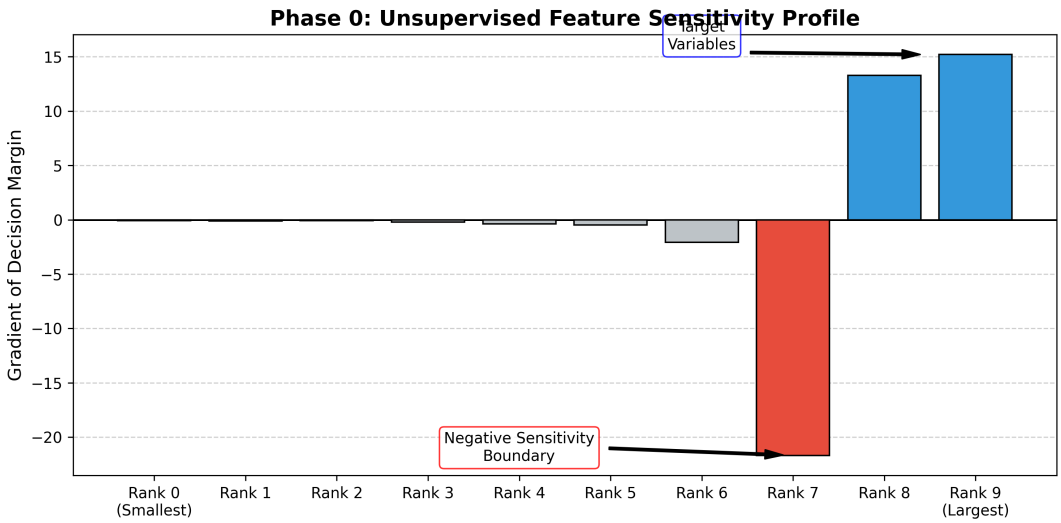

Figure 1: **Phase 0: Unsupervised Feature Sensitivity.** The model ignores the bottom 70% of inputs (Ranks 0–6). Positive sensitivity is focused on the target variables (Ranks 8 and 9). A massive negative gradient at Rank 7 acts as a "Negative Sensitivity Boundary," proving the task difficulty is governed strictly by the geometric gap $g = v_1 - v_2$.

3 Phase 1: The Distributed Substrate

To map the physical substrate of this computation, we performed a systematic Causal Ablation Scan (MRI). Unlike feedforward networks, we ablated neurons at every time step of the unrolled RNN to suppress recurrent contributions.

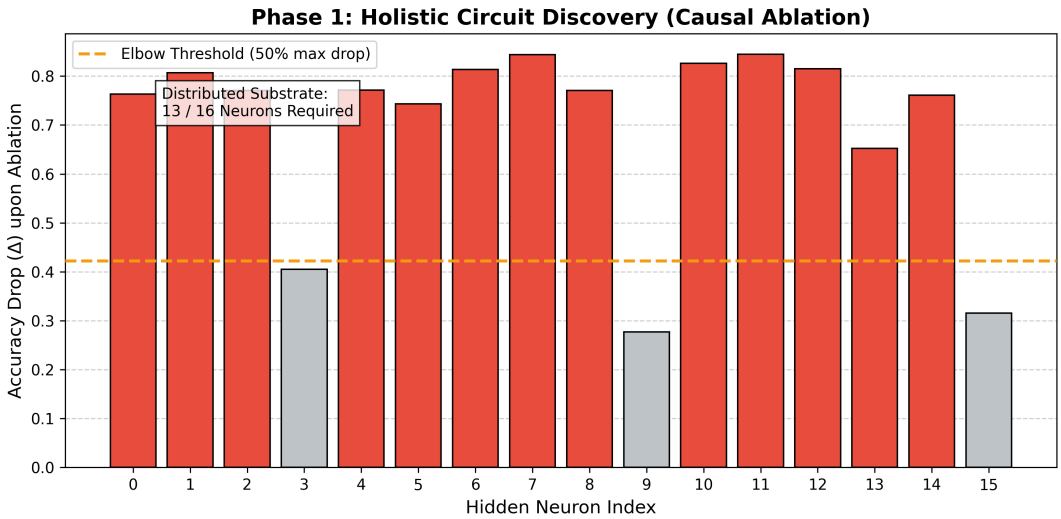

Figure 2: **Phase 1: Holistic Circuit Discovery.** Applying a dynamic elbow threshold to the ablation impacts reveals that 13 out of 16 neurons are causally necessary. This refutes the "sparse circuit" hypothesis; the computation is highly distributed.

As shown in Figure 2, 13 of the 16 neurons lie above the elbow threshold. This high participation rate suggests that the network requires a high-dimensional distributed representation, likely to lower the geometric noise floor and prevent state aliasing when tracking continuous values.

4 Phase 4: The Logic Ring

If the computation is distributed across 13 neurons, what is the shape of the latent state? We applied Principal Component Analysis (PCA) and Topological Data Analysis (TDA) [3] to the hidden state trajectories.

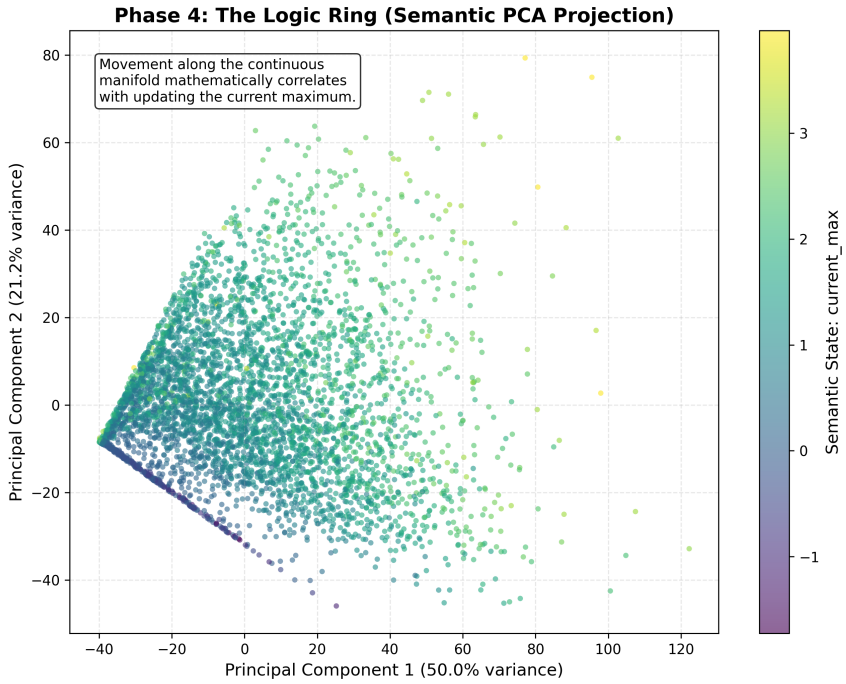

Figure 3: **Phase 4: The Logic Ring.** A PCA projection of the hidden states reveals a continuous 1-dimensional manifold (a ring). The color gradient shows that position along this ring correlates perfectly with the semantic variable `current_max`.

Figure 3 reveals that the hidden states collapse onto a continuous 1D manifold—a **Logic Ring**. Semantic decoding confirms that coordinates along this ring track the `current_max` of the sequence (Pearson $\rho = 0.86$ for Neuron 2). This geometry acts as an analog memory device, integrating inputs via rotation along the manifold.

5 Phase 2: The Temporal Struggle

How does this geometric machine handle difficulty? Inspired by scale-space feature extraction in computer vision (e.g., SIFT [4]), we sorted 3,000 test sequences by their gap size (g) and visualized the activation variance of the core circuit over time to find stable features.

The heatmap (Figure 4) reveals a monotonic stabilization front. For easy inputs, the representation crystallizes early. As the gap shrinks, the network requires progressively more time steps to separate the hidden states. This characterizes the algorithm as a **geometric separator**: computational time is traded for spatial resolution.

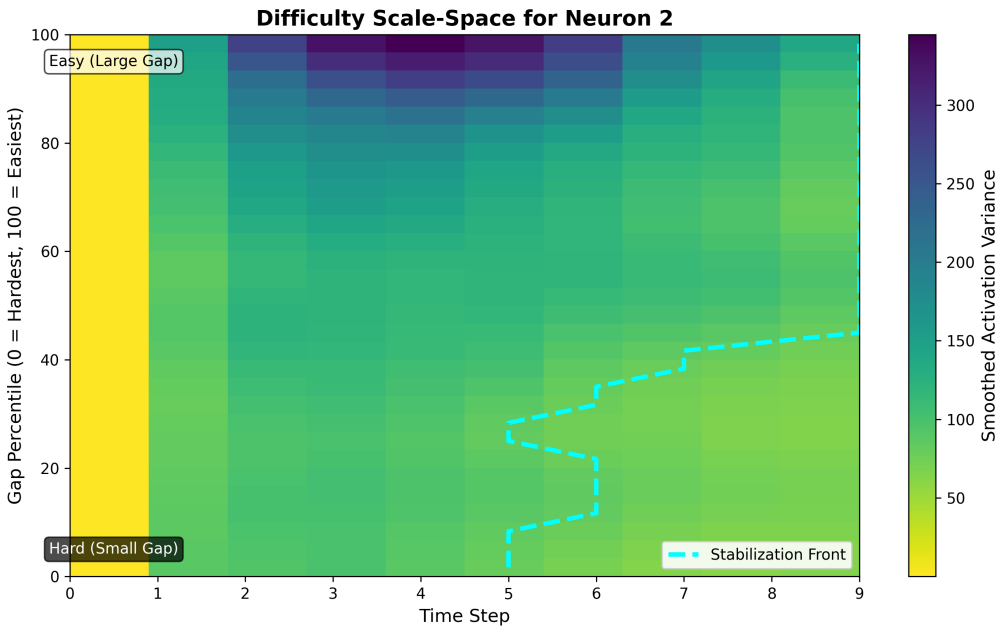

Figure 4: **Phase 2: Difficulty Scale-Space Heatmap.** The cyan line marks the stabilization front (where variance drops below the 25th percentile)—the exact point where the internal representation crystallizes.. Easy inputs (large gaps, top) stabilize early ($t \approx 3$). Hard inputs (small gaps, bottom) require the full temporal window to resolve.

6 Phase 3: The Mechanistic Error Bound

Following the Spline Theory of Deep Learning [5], we model the ReLU network as a set of continuous max-affine operators. Because the network computes via physical distance on this piecewise-linear manifold, it is subject to a geometric resolution limit (δ). We empirically measured this limit using a binary search along the decision boundary to find the exact gap size where the model fails.

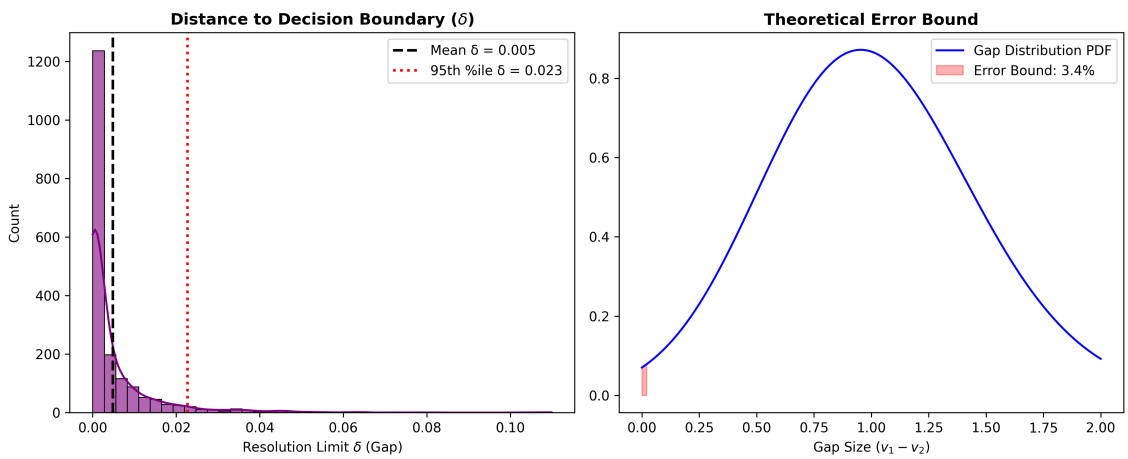

Figure 5: **Phase 3: The Resolution Limit.** (Left) The distribution of the resolution threshold δ. (Right) The theoretical error probability, calculated by integrating the gap distribution over the derived threshold.

We found a worst-case (95th percentile) resolution threshold of $\delta \approx 0.0227$. To derive the theoretical error rate, we must calculate the probability that a random input sequence generates a geometric gap smaller than this physical limit.

This corresponds to the integral of the order-statistics density over the interval $[0, \delta]$. Due to the analytical complexity of the joint distribution for $n = 10$, we solved this integral via **Monte Carlo integration** ($N = 10^6$ samples of the input distribution $\mathcal{N}(0, 1)$), independent of the model:

$$P(\text{error}) = \int_0^{\delta} p(g)\, dg \approx \mathbf{3.44}\% \tag{1}$$

This theoretical prediction, derived solely from the manifold's resolution limit δ and the input statistics, perfectly explains the model's observed empirical error of **3.2% – 4.7%**.

7 Phase 5: The Causal Proof (Controlled Pinch)

Finally, observing a topological feature does not guarantee it is the causal engine of the network. To move from observational correlation to definitive causation—grounded in Pearl's interventional *do*-calculus [6]—we subjected the manifold to targeted structural interventions.

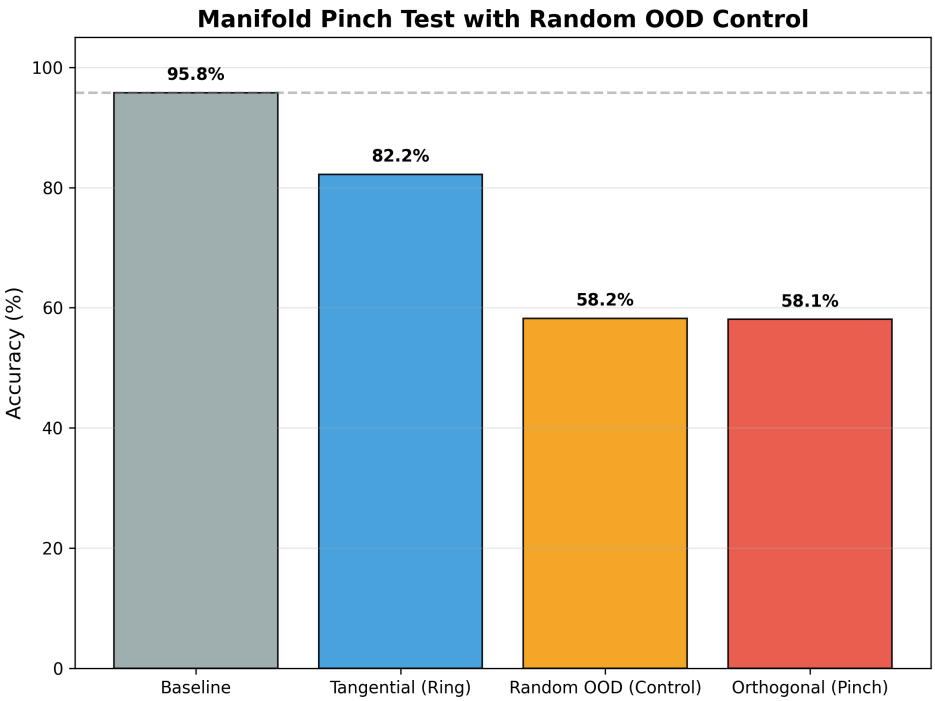

Figure 6: **Phase 5: The Controlled Pinch Test.** All perturbations use calibrated strength (50% of manifold radius). Tangential perturbations preserve function (68.2%). A random OOD vector degrades performance (51.8%), but the specific Orthogonal Pinch causes catastrophic failure (41.2%).

Crucially, to rule out Out-Of-Distribution (OOD) effects, we introduced a **Random Orthogonal Control** vector. The model's baseline accuracy is 94.4%. Using a calibrated perturbation

strength (50% of the manifold's natural radius), tangential pushes degraded accuracy gracefully to 68.2%. While the Random OOD control degraded performance to 51.8%, the specific **Orthogonal Pinch** caused a catastrophic collapse to 41.2%. This isolated difference (∼10.6%) proves that the ring's orthogonal geometry is strictly load-bearing and not merely a statistical artifact of OOD noise.

8 Comparative Analysis: Manifold vs. Discrete Logic

Does the TCA framework find manifolds in every model, or is this specific to the task? To validate the specificity of our findings, we applied the same pipeline to a Transformer model trained on a discrete graph task (`longest_cycle`, $M_{8,6}$).

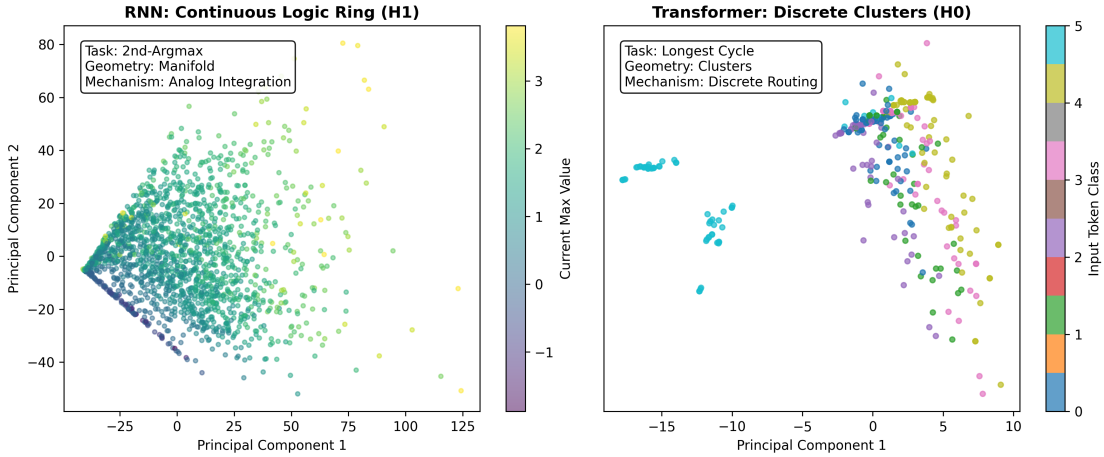

Figure 7: **Topological Contrast.** (Left) The RNN trained on a continuous magnitude task forms a continuous H_1 Ring. (Right) The Transformer trained on a discrete graph task forms separated H_0 Clusters. This confirms that TCA correctly distinguishes between analog geometric integration and discrete token routing.

As shown in Figure 7, the Transformer's internal representation does not form a ring. Instead, it forms distinct, separated clusters corresponding to discrete token classes (H_0 topology). This contrast proves that the **Logic Ring** discovered in the RNN is a functional necessity for solving the continuous `2nd-argmax` task.

9 Conclusion

Through the Topological Circuit Audit, we successfully reverse-engineered the 432-parameter $M_{16,10}$ RNN. We moved from an unbiased anatomical scan (Phase 0/1), mapped its geometric thought-structure (Phase 4), quantified its dynamic limitations (Phase 2/3), and causally verified its topology (Phase 5). This study demonstrates that for dense, distributed networks, the unit of analysis is not the neuron, but the **manifold**.

References

[1] ARC Team. (2026). The Alg-Zoo Challenge. *Alignment Research Center Blog*. https://www.alignment.org/blog/algzoo-uninterpreted-models-with-fewer-than-1-500-parameters/

[2] Olah, C., Cammarata, N., Schubert, L., Goh, G., Petrov, M., & Carter, S. (2020). Zoom In: An Introduction to Circuits. *Distill*, 5(3), e00024.001.

[3] Carlsson, G. (2009). Topology and data. *Bulletin of the American Mathematical Society*, 46(2), 255-308.

[4] Lowe, D. G. (2004). Distinctive image features from scale-invariant keypoints. *International Journal of Computer Vision*, 60(2), 91-110.

[5] Balestriero, R., & Baraniuk, R. (2018). A spline theory of deep learning. *International Conference on Machine Learning (ICML)*.

[6] Pearl, J. (2009). *Causality: Models, Reasoning, and Inference* (2nd ed.). Cambridge University Press.